(a) Base, gen. 0 training
Val. first-value acc.
1.0
0.8
0.6
0.4
0.2
0.0
50k examples
3k (pilot)
heuristic baseline
0
2000
4000
Training step
(b) Ext.: errors vs leaked mass
% of test examples / prob
6
4
2
0
inconsistent
unseen mass
incons., 50% real
0
2
4
Generation
(c) Ext.: logit attribution
Shown−unseen logit gap
6
4
2
0
total
layer-2 heads
layer-1 heads
MLPs + emb.
0
2
4
Generation
(d) Induction attention
Best induction score
1.0
0.8
0.6
0.4
0.2
0.0
base
extended
pilot (no transition)
0
2
4
Generation